

Lecture 2, R fundamentals

Most slides courtesy of Vivian Lew and Roger Peng

R is a language comprised of expressions

- ▶ Important types of expressions are:
 - ▶ Constants (e.g., 2)
 - ▶ Arithmetic expressions (e.g., $2+2$)
 - ▶ Function calls (e.g., `mean(x)` , `exp(x)`)
 - ▶ Assignments (e.g., `x <- c(1, 2, 3)`)

Recap: RStudio

- ▶ We use scripts to store R commands
- ▶ We can highlight and run a script by pressing Run or use Ctrl+Enter(Windows) , Command+Enter(Mac)
- ▶ The console echoes back and shows results
- ▶ History tab reveals the most recent submitted commands
- ▶ Help tab can be used to tell us about the R language

Everything in R is an object

- ▶ An object is storage space with an associated name.
- ▶ EVERYTHING in R is stored in an object.
- ▶ Anything that is to be used again: data, variables, functions, absolutely anything, are stored in computer memory as NAMED objects -Objects have attributes.

Example

```
test <- 2+2  
test
```

```
[1] 4
```

```
> .Last.value  
[1] 4
```

- ▶ The `.Last.value` command gets the value returned by the most recent R command.
- ▶ `test` and `.Last.value` are named objects

Everything in R is an object

`x<-3`

- ▶ We create R objects using a a “gets”. Its symbol is less than sign followed by a dash `<-` and we would read the above as “object x gets 3”
- ▶ Yes, you can use an equal sign,i.e. `x = 3`, instead of a “gets”

```
> 5  
> .Last.value  
[1] 5
```

- ▶ When we typed 5, R created an object called `.Last.value`.
- ▶ We can call it back up because we know its name.
- ▶ But typically, we use a “gets” to name and create our own R objects.

How to name objects in R?

- ▶ R is case sensitive so test is different from Test.
- ▶ R object names MUST start with a period OR a letter (and nothing else), so .Last.value is a valid R object name.
- ▶ R object names cannot have spaces (use a period or underscore instead), so .Last value is NOT a valid name.
- ▶ Periods may be found anywhere in an R name, but the underscore cannot start an R name, so __Last.value is NOT a valid name.
- ▶ No quotes around object names, so “.Last.value” is NOT a valid name.

Functions do the work in R

- ▶ Functions are a special type of R object. They exist to do something. They typically take arguments and produce a result by executing a set of operations (which are usually other functions).
- ▶ Functions have the basic form: `functionname(argument)`, so a function name, a set of parentheses and an argument. Here is a simple example

```
help(exp)
```

-But sometimes, functions don't need arguments

```
getwd()
```


Functions do the work in R

- ▶ Sometimes functions have complex arguments:
 - ▶ `plot(x <- sort(rnorm(50)),type="s",main="A plot of x, stair steps")`
 - ▶ R has a set of built-in functions available for our use.
 - ▶ However, any user can create new R functions or modify existing functions. We will cover this later.

Functions do the work in R- Example

- ▶ `exp()` computes the exponential function using Euler's number.
- ▶ Notice
 - ▶ the parentheses
 - ▶ can operate on a constant
 - ▶ can operate on other R objects
 - ▶ can be part of a larger arithmetic expression
 - ▶ can have other functions nested in them

Example

```
> exp(1) #operates on constant  
[1] 2.718282
```

```
> exp(test) #operates on the R object named test  
[1] 54.59815
```

```
> exp(1)^9 # functions can be part of arithmetic operations  
[1] 8103.084
```

```
> exp(rnorm(5)) #can also have other functions nested inside of them  
[1] 1.4237904 1.4222555 1.5741567 0.5730742 3.3567779
```

Data types in R: Vectors

- ▶ A vector is the simplest R data type.
- ▶ There are no scalars in R, a single value is a vector of length 1 (that is, `[1]`)
- ▶ R functions operate on vectors; data is basically stored as vectors or sets of vectors in R.
- ▶ There are two different kinds: atomic vectors and lists (generic vectors)
- ▶ Atomic vectors can be thought of as arrays of a single kind of value (integers, logical, characters, etc.)
- ▶ Generic vectors (AKA Lists) may contain any type of data.

Creating Vectors

-The `c()` function can be used to create vectors of objects, “c” means “concatenate”

```
x<- c(0.5, 0.6)      #numeric  
x<- c(TRUE, FALSE)   #logical  
x<- c("a", "b", "c") #character  
x<- c(2L,4L)          #integer
```

-R by default stores numbers as double. Putting capital ‘L’ after an integer forces it to be stored as an integer.

```
typeof(c(2,4))
```

```
[1] "double"
```

```
typeof(c(2L,4L))
```

```
[1] "integer"
```

Creating Vectors (Con't)

- ▶ Using colon: generates a sequence from:to in steps of 1 or -1

```
1:10  #ascending  
10:1  #descending
```

- ▶ Using seq(...): Typically wants 3 numeric values, a “from”, a “to” and a “by”.

```
> seq(1, 10, 2)  #ascending  
[1] 1 3 5 7 9  
> seq(10, 1, -2) #descending  
[1] 10 8 6 4 2
```

Creating Vectors (Con't)

- ▶ Using `vector(mode,length)` function:e.g

```
vector("logical", length = 8)
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
vector("numeric", length = 8)
```

```
## [1] 0 0 0 0 0 0 0 0
```

- ▶ Using `rep` function.

```
rep(c(0,1),5)
```

```
[1] 0 1 0 1 0 1 0 1 0 1
```

Mixing Objects

- ▶ What about the following?

```
> c(1.7, "a")  #character
[1] "1.7" "a"
> c(TRUE, 2)   # numeric
[1] 1 2
> c("a", TRUE) # character
[1] "a"  "TRUE"
```

- ▶ When different objects are mixed in a vector, coercion occurs so that every element in the vector is of the same class.
- ▶ The output vector type is determined by the highest type of the components of the vector. **logical < numeric < character.**

Explicit Coercion

- Objects can be explicitly coerced from one class to another using the `as.*` functions, if available.

```
> x<- c(0,1,2)
> typeof(x)
[1] "double"
> as.integer(x)
[1] 0 1 2
> as.logical(x)
[1] FALSE  TRUE  TRUE
> as.character(x)
[1] "0" "1" "2"
```

Explicit Coercion

- Nonsensical coercion results in NAs.

```
> x <- c("a", "b", "c")
> as.numeric(x)
[1] NA NA NA
Warning message:
NAs introduced by coercion
> as.logical(x)
[1] NA NA NA
> as.complex(x)
[1] NA NA NA
Warning message:
NAs introduced by coercion
```

Indexing - the key to unlocking data in R

- ▶ Recall that a vector in R is a numbered sequence of elements (which typically are numeric, character, or logical).
- ▶ One of the most important operations in R is accessing individual elements or subsets through indexing operations.
- ▶ Subsets of the elements of a vector may be selected by appending to the name of the vector an index vector in square brackets.
- ▶ The format is `vector1[vector2]`, with the result that we select those elements of `vector1` whose indices are given in `vector2`.
Vector2 is an index vector

Example

```
> rivers #built-in vector river
[1] 735 320 325 392 524 450 1459 135 465 600 330 336 280 315 870
[16] 906 202 329 290 1000 600 505 1450 840 1243 890 350 407 286 280
[31] 525 720 390 250 327 230 265 850 210 630 260 230 360 730 600
[46] 306 390 420 291 710 340 217 281 352 259 250 470 680 570 350
[61] 300 560 900 625 332 2348 1171 3710 2315 2533 780 280 410 460 260
[76] 255 431 350 760 618 338 981 1306 500 696 605 250 411 1054 735
[91] 233 435 490 310 460 383 375 1270 545 445 1885 380 300 380 377
[106] 425 276 210 800 420 350 360 538 1100 1205 314 237 610 360 540
[121] 1038 424 310 300 444 301 268 620 215 652 900 525 246 360 529
[136] 500 720 270 430 671 1770

> rivers[28]
[1] 407
> rivers[132:137]
[1] 525 246 360 529 500 720
> rivers[c(1,3,5,7)]
[1] 735 325 524 1459
rivers[seq(5,120,10)]
[1] 524 870 1243 327 600 259 332 260 696 460 377 1205
> rivers[rivers > 2000]
[1] 2348 3710 2315 2533
> which(rivers>2000) #This is different, understand the difference
[1] 66 68 69 70
```

Data types in R: Matrices

- ▶ Matrices are vectors with a dimension attribute. The dimension attribute is itself an integer vector of length 2 (nrow, ncol).i.e. the number of rows and the number of columns.
- ▶ A matrix, like an atomic vector, may only have one type of data (e.g., all numeric, all character, all logical)
- ▶ The rows of a matrix must all be the same size
- ▶ The columns of a matrix must all be the same size
- ▶ But the size of the rows are not necessarily equal to the size of the columns
- ▶ And like vectors, there are many ways to create matrices

Matrices (Con't)

- ▶ Matrices are filled *column-wise*, so entries can be thought of starting in the “upper left” corner and running down the columns.

```
> m <- matrix(1:6, nrow = 2, ncol = 3)
> m
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

- ▶ If we want to fill a matrix *row-wise*, we need to set `byrow = TRUE`. The default is `FALSE`.

```
> m <- matrix(1:6, nrow = 2, ncol = 3, byrow = TRUE)
> m
```

	[,1]	[,2]	[,3]
[1,]	1	2	3
[2,]	4	5	6

Matrices (Con't)

- ▶ Matrices can also be created directly from vectors by adding a dimension attribute.

```
> m<- 1:10
> m
[1] 1 2 3 4 5 6 7 8 9 10
> dim(m) <- c(2, 5)
> m
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    3    5    7    9
[2,]    2    4    6    8   10
```

Matrices (Con't)

- ▶ Matrices can be created by column-binding or row-binding with `cbind()` and `rbind()`.

```
> x<-1:3
> y<-10:12
> cbind(x, y)
      x  y
[1,]  1 10
[2,]  2 11
[3,]  3 12
> rbind(x, y)
  [,1] [,2] [,3]
x    1    2    3
y   10   11   12
```


Examples

```
> a<- 1:15
> # matrix function needs either the number of rows or
> #the number of columns
> matrix(a,nrow=5)
      [,1] [,2] [,3]
[1,]     1     6    11
[2,]     2     7    12
[3,]     3     8    13
[4,]     4     9    14
[5,]     5    10    15

> matrix(a,ncol=5)
      [,1] [,2] [,3] [,4] [,5]
[1,]     1     4     7    10    13
[2,]     2     5     8    11    14
[3,]     3     6     9    12    15

> matrix(c("a","A","b","B"), nrow=2, byrow=TRUE)
      [,1] [,2]
[1,]  "a"  "A"
[2,]  "b"  "B"

> matrix(c(TRUE,TRUE,FALSE,TRUE), nrow=2)
      [,1] [,2]
[1,]  TRUE FALSE
[2,]  TRUE  TRUE
```

Example: Creating a matrix and indexing

```
> BOD2<- as.matrix(BOD) #as.matrix forces BOD to become a matrix
> BOD2
      Time demand
[1,]    1    8.3
[2,]    2   10.3
[3,]    3   19.0
[4,]    4   16.0
[5,]    5   15.6
[6,]    7   19.8
> dim(BOD2)
[1] 6 2
> length(BOD2)
[1] 12
> nrow(BOD2)
[1] 6
> ncol(BOD2)
[1] 2
> BOD2[9] #If we wish to extract a single cell from matrix
[1] 19
> BOD2[3,2] #clearer
demand
      19
> # If we wish to extract a specific row
> BOD2[3,] # gives row 3, all columns
      Time demand
      3    19
> # If we wish to extract a specific column
> BOD2[,2]
[1] 8.3 10.3 19.0 16.0 15.6 19.8
> BOD2[4:6, 1:2] #rows 4,5,6 of columns 1 and 2 (here same as BOD2[4:6,])
      Time demand
[1,]    4   16.0
[2,]    5   15.6
[3,]    7   19.8
#Try BOD2[c(4,6), 1:2]
```

Data types in R: Lists

- ▶ Lists are a special type of vector (generic vector) that can contain elements of different classes.
- ▶ A list allows you to gather a variety of (possibly unrelated) objects under one name.

```
> mylist<- list(c(2,3,4), c(TRUE,FALSE), c("A","B","C"))
> mylist
[[1]]
[1] 2 3 4

[[2]]
[1] TRUE FALSE

[[3]]
[1] "A" "B" "C"
```

Identify elements of a list using the `[[]]` convention.

```
> mylist[[1]] # first component of the list
[1] 2 3 4
> mylist[[3]] #3rd element of the list
[1] "A" "B" "C"
> mylist<- list(Jane = c(2,3,4), Dave = c(TRUE,FALSE), Mary = c("A","B","C"))
> mylist
$Jane
[1] 2 3 4

$Dave
[1] TRUE FALSE

$Mary
[1] "A" "B" "C"

> mylist$Jane
[1] 2 3 4
```

Data types in R: Data Frames

- ▶ A data frame is a rectangular list.
- ▶ In other words, they are represented as a special type of list where every element of the list has to have the same length.
- ▶ A data frame is like a matrix with columns possibly of differing modes and attributes (e.g., numeric columns, character columns, logical columns).
- ▶ Data frames typically have “observations” in the rows and “variables” in columns in a rectangular matrix-like structure (equal length rows and equal length columns).

Data Frames - Example

```
> x <- data.frame(foo = 1:4, bar = c(T, T, F, F))
> x
  foo  bar
1   1 TRUE
2   2 TRUE
3   3 FALSE
4   4 FALSE
> nrow(x)
[1] 4
> ncol(x)
[1] 2
```

Data frames (Con't)

- ▶ Data frames have a `row.names()` attribute which labels the observations.
- ▶ Data frames have a `names()` attribute which labels the variables.
- ▶ Data frame rows and columns are accessed using matrix indexing rules OR list indexing rules.

Data frames -Example

```
> x
  foo  bar
1   1 TRUE
2   2 TRUE
3   3 FALSE
4   4 FALSE
> row.names(x)
[1] "1" "2" "3" "4"
> names(x)
[1] "foo" "bar"
> names(x) <- c("Dave" , "Jane")
> x
  Dave  Jane
1    1 TRUE
2    2 TRUE
3    3 FALSE
4    4 FALSE
> x[1,2]
[1] TRUE
> x[,2] # like a matrix extracts all rows, column 2
[1] TRUE TRUE FALSE FALSE
> x[[1]] #extracts a vector from a list
[1] 1 2 3 4
> x$Jane # extracts a named column
[1] TRUE TRUE FALSE FALSE
```


Swirl Package

- ▶ For some of the problem sets we'll be using the swirl software package for R in order to illustrate some key concepts. The swirl package turns the R console into an interactive learning environment.
 - ▶ Install swirl using: **install.packages("swirl")**
 - ▶ Load Swirl using: **library(swirl)**
 - ▶ Install the R Programming course using:
install_from_swirl("R Programming")
 - ▶ Start swirl and complete the lessons using: **swirl()**